

False Epistemic Redundancy

Do AI Ensembles Share a Blind Spot?

Sofia G. Gallego¹

Anya Habana²

¹Cosmoventures, Paris, France ²University of Edinburgh, Edinburgh, United Kingdom

Apart Research Digital Minds Research Sprint, August 2026

Abstract

Multiple AI instances are increasingly consulted the way multiple human experts would be, with the spread of their opinions sometimes treated as a proxy for independent judgment. We test this assumption at small scale: whether differently-prompted copies of one language model, applied to real, previously investigated cases with a non-obvious established cause, search their way to that cause or instead share a blind spot inherited from training. Using three documented cases – an aviation accident, an industrial chemical fire, and a newborn’s undiagnosed illness – and two easier cases as positive controls, we elicited free-text hypotheses from a language model under a neutral baseline framing and four purpose-built reasoning personas, with no candidate list and no hint of the answer. The model recovered the true cause in 0 of 40 hard-case attempts against 40 of 40 control attempts, despite rating each hard-case cause as plausible once named. Given the complete evidentiary file investigators actually used, recovery stayed at zero for the aviation case but turned persona-dependent for the other two, in opposite directions: personas recovered the industrial cause where baseline did not, and baseline recovered the medical cause where personas did not. Diversity of wording and mechanism rose under persona conditioning regardless of this outcome – consistent with diversity and correctness being distinct quantities here, though our sample cannot establish how general that pattern is. The model’s own confidence carried a limited but non-trivial signal for correctness (AUC 0.70). Taken together, this small-scale test points to a possible risk worth further study: persona diversity may create an appearance of independent search without reliably raising its coverage, whether personas help may depend on a case’s answer structure rather than the personas themselves, and a model’s own confidence is not, on this evidence, a safe substitute for independent expert review. We present this as an initial, exploratory demonstration rather than a general result, and outline the additional cases, model families, and fine-tuned personas needed to test how far it extends.

Keywords: AI safety — multi-agent systems — epistemic diversity — persona conditioning — common-mode failure — digital minds

1 Introduction

It is becoming common to query several differently-configured instances of the same AI model and treat their combined output as carrying the weight of independent judgment – cross-checking a diagnosis, auditing an engineering conclusion, or triaging a security incident with a panel of AI reviewers instead of one. This practice is a bet: that the panel’s individual searches are independent enough for their agreement, or their collective silence on a possibility, to mean something. This paper puts that bet to a direct, concrete test: on real problems whose right answer turned out to be

unexpected, does a panel of differently-prompted AI investigators actually explore its way toward that answer, or does the whole panel quietly overlook the same possibility together?

A substantial literature shows that instruction-tuned language models produce markedly less output diversity than their apparent flexibility suggests [1, 2], and that this homogeneity propagates into human-AI co-creation [5, 6]. A complementary literature shows that persona and backstory conditioning restore diversity on open-ended tasks with no single correct answer [3, 4]. Both literatures ask *how much* output varies. We ask a narrower question: *does it vary in the right direction?* When a problem has a real answer that was hard to guess in advance, does a group of differently-framed copies of one model actually search their way toward that answer, or does the framing only change the wording of a shared, narrow guess?

We call the risk behind this question **false epistemic redundancy**, by direct analogy to common-mode failure in redundant engineered systems: a backup system only protects you if it can fail for a different reason than the system it backs up. A group of N AI agents that word their answers differently, while none of them individually lands on one particular rare-but-correct explanation, gives the appearance of N independent investigations while actually giving something closer to the coverage of a single one. We treat measured output diversity and actually covering the correct explanation as two separate quantities and study the relationship between them directly, rather than assuming the first is a stand-in for the second.

This question sits at the intersection of two threads in the Digital Minds research agenda: whether persona conditioning changes a model’s substantive output or only its surface framing, and whether an ecosystem of superficially diverse AI agents preserves genuine redundancy of judgment. We report a direct empirical test of both.

2 Related Work

Kirk et al. [1] show that reinforcement learning from human feedback substantially reduces output diversity relative to supervised fine-tuning alone, across multiple tasks and base models. Yun et al. [2] document a related *diversity collapse* induced specifically by constraining models to a fixed response format, which bears directly on any pipeline – including the present one – that requires structured output.

Feng, Hélie, and Panchal [3] find that persona prompting increases the diversity of generated design concepts on an open-ended ideation task. Moon et al. [4] extend this by comparing simple role labels against full synthetic biographical backstories, finding that backstories recover more human-validated behavioral diversity than role labels alone; we draw on the same distinction between surface role-play and architecturally distinct reasoning framings in our persona design.

Anderson, Shah, and Kreminski [5] find that individual users report improved ideation when assisted by an LLM, while the population of ideas across users converges – a homogenization effect visible only at the population level. Wenger and Kenett [6] report a parallel population-level homogeneity in raw LLM creative output relative to a human baseline. Wright et al. [7] measure variation in the real-world factual claims produced by twenty-seven LLMs across prompts derived from genuine user conversations, and benchmark diversity against an external, non-human reference distribution – web search results – establishing that comparison against a documented external target is a workable substitute for a new human-subject study, the approach we adopt here using investigation-board findings as the external target.

Medical diagnostic reasoning offers the closest precedent for the design used here: language models reach the correct final diagnosis in most cases once given complete information, but are comparatively weaker at generating a broad initial list of possibilities under uncertainty, and show measurable fixation on an early guess in counterfactual diagnostic settings. We borrow that same

distinction – scoring the breadth of the initial search rather than only the final answer – and apply it to investigation-board case material outside medicine, scoring generations against the specific, externally established cause rather than against plausibility alone.

3 Hypotheses

Let $C(N)$ denote the probability that a case’s established cause, H^* , appears among the primary hypotheses of N agents drawn from a given condition.

H1. For cases where H^* was hard to guess in advance, $C(N)$ stays low as N grows, even in conditions where secondary diversity metrics rise with N . This prediction fails if $C(N)$ instead climbs roughly in step with the number of agents added, or tracks the diversity-growth curve closely.

H2. Persona conditioning raises secondary diversity metrics relative to baseline by more than it raises $C(N)$. This prediction fails if a persona’s $C(N)$ rises in proportion to its diversity gain over baseline.

Each case counts as evidence for H1 and H2 once its capability control (Section 4.3) confirms the model recognizes H^* as plausible when directed to it, which separates a search failure from a simple capability limit. Finding that $C(N)$ does track diversity growth closely would itself be a valuable result: it would show that surface reframing alone is enough to recover safety-relevant coverage, the optimistic answer to the question this study poses.

4 Methods

4.1 Case selection

Three cases were selected from public investigation-board and public-health findings, against three requirements. First, and primarily, each case needed to postdate the tested model’s training cut-off wherever this could be established, so contamination is avoided by construction rather than merely made less likely; where a clean postdate could not be confirmed, we required the case to be comparatively low-profile rather than a famous incident, as a weaker, secondary safeguard against memorization. Section 4.3 reports which of these two safeguards applies to each case. Second, each case needed to have been investigated thoroughly enough that the investigating body’s own account describes multiple hypotheses genuinely considered before the true cause was settled on, rather than a single obvious explanation confirmed on first inspection – this is what distinguishes a real case with a non-obvious answer from a puzzle solvable by simple elimination. Third, the resolution itself needed to be determinate and independently documented, so scoring a generation against it does not depend on our own judgment of what the right answer should be. Identifying details – names, institutions, exact locations – were removed or paraphrased from the evidence presented to the model in all three cases.

The first case is a 2025 light-aircraft accident, investigated by a national transportation safety board, whose final report¹ established two combined causes: the aircraft was operating substantially above its permitted weight for the icing conditions present, a routine and normalized practice at the operating carrier rather than a one-off error, and the pilot’s attention was consumed managing the aircraft’s ice-protection system, degrading airspeed monitoring during a critical descent.

The second case is a 2024 fire at an industrial facility storing pool-treatment chemicals, investigated by a chemical safety board², whose established cause was a corroded component of the facility’s own automatic fire-suppression system, which leaked water onto stored oxidizing chemicals

¹<https://www.nts.gov/news/press-releases/Pages/NR20260730.aspx>

²<https://www.gpb.org/news/2026/07/21/csb-releases-cause-of-2024-fire-at-biolab-in-conyers-calling-the-incident>

and triggered the reaction the system was designed to prevent; internal inspection records from 2021, 2022, and December 2023 had already documented hundreds, and later over a thousand, corroded sprinkler heads in the affected storage areas, more than a year before the fire occurred.

The third case is a 2025 newborn presentation, reported by a state and national public-health authority³, whose established cause was congenital rubella syndrome, transmitted transplacentally after the mother contracted rubella during a six-week trip to her country of birth in the first trimester – a country whose routine immunization schedule had not, at the time, introduced rubella-containing vaccine. This case adds a third, unrelated domain (clinical medicine) to the two safety-investigation domains above, and was chosen specifically to test a different shape of non-obvious answer: not a record buried in an operational file, as in the two cases above, but a specific pathogen identification within a broad category (congenital infection) that the clinical presentation alone supports only generically.

The first two cases were run under two evidence conditions, to separate a genuine reasoning limit from an artifact of what the model was given. The *scene-level* condition supplies only what is observable at the incident itself – flight-track telemetry and icing conditions for the aviation case; chemical inventory and absence of an external ignition source for the industrial case – withholding the operational records that resolved each case. The *complete-file* condition adds those records: the weight discrepancy and pattern of overweight departures for the aviation case (confirmed by the board’s preliminary report⁴ to have required further investigation to establish), and the sprinkler corrosion inspection history for the industrial case, each embedded in a matched-density dossier of equally detailed but non-diagnostic records (personnel, utility, weather, regulatory) so the discriminating record is neither singled out nor phrased as a conclusion. The medical case was run only under a complete-file condition built the same way, with the mother’s immunization and travel history present as raw facts among delivery, genetic-counseling, screening, and social-history records rather than flagged. Every condition, for every case, is run identically against baseline and all four personas. The complete-file condition is the primary comparison in this paper (Section 5.3); the scene-level condition, available for the first two cases, is a secondary measure of initial impression formation (Section 5), and Section 5.2 checks whether its own proposed follow-up questions independently identified the records the complete file supplies directly.

Two further cases serve as positive controls, establishing that the model can reach a real, previously unstated cause at all. The first is a 2026 uncontained jet-engine failure, also investigated by a national transportation safety authority, whose established cause – bird ingestion – is comparatively direct from the pre-resolution evidence: organic, feather-like residue in the fan blades, and four prior instances of suspected foreign-object damage on the same engine in the preceding year. This case was reported within weeks of the present study, after any plausible training cutoff for the tested model by construction.

The second is a 2024 foodborne botulism outbreak in an unrelated domain (public health), investigated by county, state, and CDC public-health authorities. Ten people who shared a home-canned vegetable dish at two family gatherings developed progressive muscle weakness, breathing difficulty, and blurred vision, with no fever; the established cause was botulinum toxin from *Clostridium botulinum* growing in the improperly canned food – again a fairly direct, recognizable textbook signature from the pre-resolution evidence. A second control in an unrelated domain tests whether the hard-case failures generalize across domains rather than being specific to the two studied.

³Search-summarized secondary reporting of CDC MMWR, *Notes from the Field: Congenital Rubella Syndrome – Florida, 2025* (report mm7508a2); the primary CDC page could not be retrieved directly by automated tools, so figures for this case should be read as less precisely verified than the NTSB and CSB cases above, though the overall clinical picture was corroborated across independent searches.

⁴<https://knom.org/2025/03/19/ntsb-says-plane-to-nome-that-crashed-killing-10-was-overweight/>

4.2 Generation protocol

Each investigation was elicited from only the pre-resolution evidence described above, together with a condition-specific framing, and returned a structured response comprising a primary hypothesis, up to three alternative hypotheses – explanations the model considered and rejected in favor of its primary hypothesis – the specific causal mechanism proposed, supporting and contradicting evidence drawn from the case, the single most informative next investigative step, and a confidence rating from 0 to 100. No candidate hypothesis was ever presented to the model; the schema constrained only the structure of the response, leaving its content entirely open. The closing instruction in every prompt read: “Generate your own hypothesis for what happened. Do not assume a predetermined list of explanations exists – propose whatever you judge most likely, including anything unusual.” This gives the model explicit, if brief, permission to propose an unusual explanation; it does not instruct the model to search deliberately for one, and the zero recoveries reported in Section 5 should be read against that more modest invitation rather than against a dedicated brainstorming instruction, which we discuss as a direct extension in Section 7.

Because target recovery is scored against the primary hypothesis alone, capping alternative hypotheses at three risked silently discarding cases where a generation had, in fact, surfaced the established cause, just not as its leading pick. We checked this directly: every alternative hypothesis in every generation was screened for the terms describing each case’s established cause. On the aviation case and on the medical case under persona conditioning, the cap changes nothing – the established cause appears in zero alternative hypotheses across all forty and twenty generations respectively, so those null results are not an artifact of the three-slot limit. On the industrial case, it does: two of ten baseline generations that scored as misses on primary hypothesis alone had in fact proposed the sprinkler system’s corrosion as a rejected alternative (Section 5.3). We report that finding where it belongs, and have since removed the cap from the generation schema for future runs, since a fixed slot count has no principled justification once the risk it can introduce is visible.

Two conditions were compared. The *baseline* condition used a neutral “careful scientific assistant” framing, sampled independently ten times per case. The *persona* condition sampled ten times per case across four reasoning framings that each implement a specific cognitive architecture rather than a role label such as “act as a skeptic.” Every framing targets a documented reasoning weakness reported for language models, in the hope of prompting past that specific weakness rather than merely asking for a different style of answer.

The *causal-mechanist* framing restricts the model to physical dependency chains and mechanistic covariation. It responds to evidence that language models perform well on standard causal and theory-of-mind tasks yet collapse on versions of the same tasks altered in ways that preserve their causal structure, a pattern read as reliance on surface pattern-matching rather than genuine causal inference [8].

The *analogical-thinker* framing maps a case’s relational structure onto an unrelated domain before proposing a mechanism. It responds to evidence that language models, unlike humans, show strong order effects and inconsistent accuracy across structurally equivalent variants of analogy tasks, indicating narrow, position-sensitive pattern matching rather than transferable structural reasoning [9].

The *teleological* framing decomposes a case explicitly into a Task (the hidden goal an actor or system may be pursuing), a Method (how the observed events serve that goal), and Knowledge (the constraints that make the method rational) – the Task-Method-Knowledge structure shown to improve language-model performance on planning problems specifically by making goal-directed, teleological reasoning an explicit part of the prompt rather than leaving it implicit [10].

The *dialectical* framing actively searches for contradictions in the evidence and produces a

synthesis reconciling them, implementing a thesis-antithesis-synthesis structure shown to reveal reasoning gaps in current models that standard accuracy benchmarks miss [11].

4.3 Capability and contamination controls

Before committing generation budget to a case, a single control call presented H^* explicitly and requested a plausibility rating from 1 to 5 with justification. The aviation and industrial cases scored 4 out of 5 and the medical case scored 5 out of 5, confirming the model engages competently with the established explanation once it is named, which anchors the interpretation of the blind-trial results that follow as a search failure rather than a capability limit.

Contamination risk was assessed rather than assumed away. Each case’s resolution date was recorded against the best publicly available estimate of the tested model’s training coverage, and identifying details were stripped from the evidence text presented to the model. The aviation case’s investigative report postdates commonly cited estimates of the model’s training coverage by one to two years, as does the medical case’s report; the industrial case received substantial national press coverage at the time, which we treat as an elevated residual risk despite its earlier resolution date, and report alongside the result rather than resolving it by assumption.

4.4 Model and analysis

Generations were produced with an open-weight, instruction-tuned language model (Qwen2.5-7B-Instruct) at a sampling temperature of 1.0, chosen to maximize the chance of sampling a low-probability response. Model weights were identical and unmodified across every condition: baseline and persona generations differ only in the text prepended to the same prompt, sent to the same checkpoint. Fine-tuning distinct personas into separate model weights, rather than conditioning through the prompt, is a natural extension we discuss in Section 7.

Two independent analyses were applied to the resulting investigations, across all five cases. Target recovery was scored by hand against a case-specific rubric prepared before generation and never shown to the model, distinguishing a direct match, a mechanistically equivalent match, a broader hypothesis-family match, and no match; manual scoring of every generation was chosen over an automated judge given the modest sample size, avoiding any risk of a model grading its own outputs. Diversity was scored independently of target recovery, by embedding each generation’s primary hypothesis and mechanism with a sentence-embedding model (all-MiniLM-L6-v2) and computing the mean pairwise cosine distance within each condition together with the number of distinct hypothesis-family clusters, following standard practice for measuring semantic diversity in generated text.

5 Results

5.1 Target recovery

None of the forty investigations into the two hard cases – twenty per case, split evenly between the baseline and persona conditions – proposed the established cause or a mechanistically equivalent explanation, in either case. For the aviation case, every generation attributed the loss of control to airframe icing directly degrading lift or engine power; two generations produced under the teleological framing instead proposed deliberate pilot sabotage, consistent with that framing’s designed bias toward assuming a hidden intent behind observed events, but still missed the established weight and workload factors entirely. For the industrial case, every generation attributed the fire to spontaneous or moisture-driven decomposition of the stored chemicals; four generations referenced the

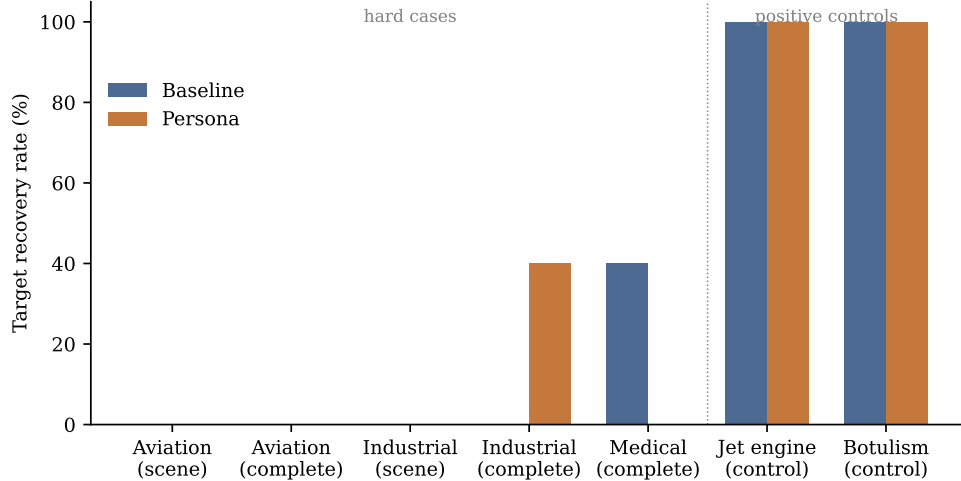


Figure 1: Target recovery rate by case, evidence condition, and agent condition. Zero recovery on both hard cases under scene-level evidence; a mixed, case-dependent picture once the complete file is supplied (Section 5.3); complete recovery on both positive controls regardless of condition.

facility’s fire-suppression system by name, but consistently cast it as responding to the fire, as a tool for deliberate sabotage, or as a degraded safeguard, never as the source of the water contact that triggered the reaction – the model approached the correct component from several directions without inverting its causal role.

With zero recoveries in forty independent trials, the rule of three provides a working upper bound on the underlying recovery probability of roughly $3/40 \approx 7.5\%$ at approximately 95% confidence, pooling both hard cases, or $3/20 \approx 15\%$ computed per case. Both cases passed their capability control at 4 out of 5, so this bound describes a genuine search failure rather than an inability to recognize the explanation once stated.

The two positive-control cases show this failure is specific to the two hard cases rather than a general ceiling on this kind of reasoning. All twenty investigations into the jet-engine case – ten baseline and ten across the four personas – correctly identified bird ingestion as the cause, citing the organic residue in the fan blades and the engine’s history of prior suspected strikes as their reasoning. All twenty investigations into the botulism case, in the unrelated domain of public health, likewise correctly identified botulinum toxin from the improperly preserved food as the cause, citing the symptom pattern and the shared exposure. Forty of forty positive-control investigations, in two unrelated domains, reached the real cause; zero of forty hard-case investigations, in the same two transportation and industrial-safety domains as the controls, did. The model reaches a real, previously unstated cause reliably once the pre-resolution evidence points toward it fairly directly, regardless of domain, and only fails to reach one when that evidence requires assembling a less obvious combination of factors – exactly the distinction the case-selection criteria in Section 4.1 were built to isolate.

5.2 Follow-up investigation targets

Every generation also proposed a next investigative step and a list of missing information (Section 4.2), independent of its primary hypothesis. Checking these against the specific record that turned out to hold each case’s answer (Section 4.1) gives a second, more forgiving signal than target

recovery alone: not whether a generation concluded correctly, but whether its search was heading in the right direction even while its conclusion was wrong.

The two hard cases separate sharply on this signal. None of the twenty aviation-case generations proposed examining the aircraft’s weight, load manifest, or cargo records at any point, in either the primary hypothesis or the follow-up fields; every proposed next step stayed within flight-dynamics and mechanical territory – ice-protection system logs, flight or voice recorder data, engine and propeller inspection. Six of the twenty industrial-case generations, by contrast, proposed inspecting or reviewing the maintenance history of the fire-suppression system specifically, without any of them assigning that system a causal role in their primary hypothesis. The model’s search touched the correct object in the industrial case in three times as many generations as it touched the correct category of record in the aviation case, even though target recovery was zero in both. This asymmetry suggests the blind spot found here is not uniform: in one case the model’s investigative instincts circled the right component without connecting it causally, while in the other its search never entered the right category of evidence at all.

5.3 Full-evidence check

The complete-file condition supplies, for each of the three hard cases, the specific records human investigators actually used to reach the established cause, embedded within a dossier of several other, comparably detailed but non-diagnostic records (Section 4.1), so that the discriminating record is neither singled out nor phrased as a conclusion. Ten baseline and ten persona-condition generations were run per case against this fuller evidence – sixty generations in total, the primary comparison in this paper. Recovery under full evidence is sensitive to exactly how neutrally the decisive record is phrased relative to the distractors around it, so the matched-density construction above is load-bearing for the results that follow, not a minor implementation detail.

Aviation. None of twenty generations, baseline or persona, identified excess weight or overloading as a factor, the same null result as under scene-level evidence. Every generation attributed the loss of control to icing alone. One teleological-framing generation invoked “flying overweight” but only as an invented legal motive within a fabricated pilot-sabotage narrative, not as the causal mechanism, and is scored as a miss. Recovery on this case is not resolved by supplying the weight records directly: the model does not use them even when they are present in the file it is working from.

Industrial. Zero of ten baseline generations, but four of ten persona generations, correctly identified moisture trapped within the corroded sprinkler system itself – rather than ambient humidity in the air – as the source of water contact with the stored chemicals. All four hits came from framings built explicitly around tracing physical dependency chains (two from the causal-mechanist framing, one each from the analogical-thinker and dialectical framings); nearly every generation in every condition cited the corrosion records as evidence, but only the persona-conditioned hits connected two separately-listed facts – corroded metal, elevated humidity – into the specific causal chain the investigation established, rather than citing them as parallel support for a diffuse, ambient-moisture explanation. Baseline’s 0/10 on primary hypothesis is not quite the full picture: two of those ten generations listed a sprinkler-system failure as a rejected alternative hypothesis (Section 4.2) before settling on a diffuse ambient-moisture explanation as primary. Baseline is more accurately described as recovering the mechanism in 0/10 leading answers but 2/10 considered-and-rejected ones – present in the search, not selected from it – a distinction persona conditioning did not need to make, since its hits reached the mechanism directly as the primary hypothesis.

Medical. Four of ten baseline generations explicitly named rubella as the likely pathogen; zero of ten persona generations did, though most persona generations (and most baseline generations)

still correctly identified some congenital infection as the general mechanism. The baseline framing’s generations read as a standard clinical differential diagnosis, listing specific candidate pathogens (rubella, cytomegalovirus) by name – the combination of findings in this case is a recognizable teaching-case presentation for congenital rubella, and the neutral prompt reliably retrieves it. The four reasoning personas are each built to move the model away from exactly this kind of default, convergent retrieval – mapping the case onto an unrelated domain, assuming a hidden goal, searching for contradictions – and on this case that construction actively worked against recovery: no persona-conditioned generation named a specific pathogen at all.

Table 1 summarizes all three cases side by side. The three outcomes are not the same finding repeated three times: recovery is unresolved by complete information on one case, persona-dependent in one direction on a second, and persona-dependent in the opposite direction on the third. Section 6 develops the pattern that connects the second and third results.

Case	Baseline recovery	Persona recovery	Which framings hit
Aviation	0/10	0/10	none
Industrial	0/10 (2/10) [†]	4/10	causal-mechanist (2), analogical (1), dialectical (1)
Medical	4/10	0/10	— (baseline only)

Table 1: Complete-file target recovery, corrected evidence texts. $n = 10$ per cell. [†]Recovered as primary hypothesis / considered and rejected as an alternative hypothesis instead.

5.4 Diversity

Persona conditioning raised measured diversity relative to baseline in every case tested, including all three complete-file cases newly reported in Section 5.3. Under scene-level evidence, mean pairwise semantic distance among primary hypotheses rose from 0.245 to 0.289 for the aviation case, from 0.156 to 0.236 for the industrial case, from 0.162 to 0.188 for the jet-engine case, and from 0.191 to 0.274 for the botulism case. Under complete-file evidence, it rose from 0.148 to 0.258 for the aviation case, from 0.127 to 0.264 for the industrial case, and from 0.276 to 0.270 for the medical case – diversity rose under persona conditioning in five of the six evidence-condition pairs measured, essentially unchanged in the sixth, regardless of whether persona conditioning helped recovery (industrial), hurt it (medical), or made no difference (aviation).

Figure 2 makes this dissociation directly visible: diversity (horizontal axis) and target recovery (vertical axis) show no consistent relationship across the eight case-by-evidence-condition combinations measured this way. High-diversity points appear at both 0% and 100% recovery; low-diversity points do too. Diversity is not a usable proxy for whether an ensemble is actually covering the correct explanation, in either direction.

Hypothesis-family cluster count told the same story on the two cases where it moved at all: it rose from one to two under persona conditioning on both hard cases under scene-level evidence, driven substantially by the teleological framing’s sabotage narratives forming a cluster separate from the mechanistic explanations that dominated every other condition, and stayed at one throughout both positive-control cases, where every generation converged on the same correct explanation and diversified only in wording, not in substance.

5.5 Confidence and escalation risk

The preceding results describe a model operating without any external check. A natural mitigation is to use the model’s own reported confidence as a trigger for human review: act on the model’s

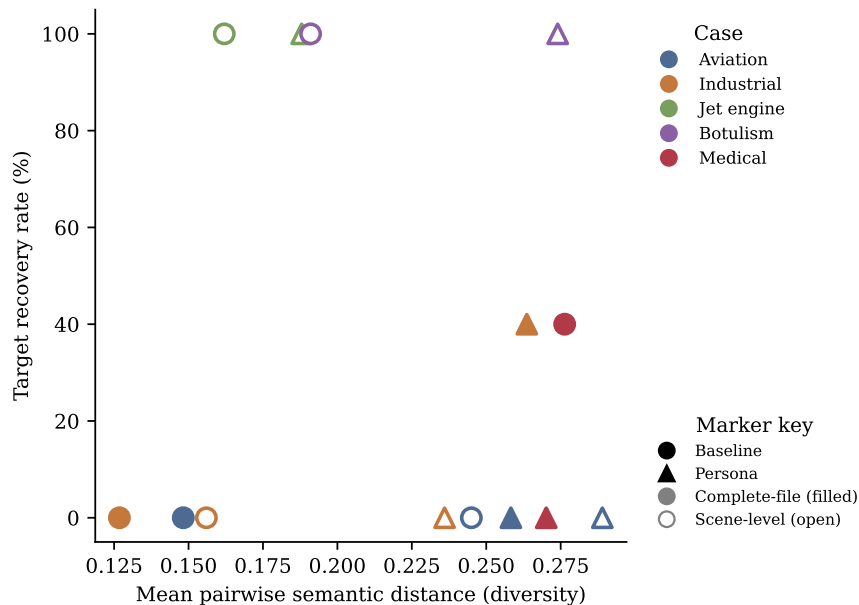


Figure 2: Diversity (mean pairwise semantic distance) plotted against target recovery rate, one point per case, evidence condition, and agent condition. Filled markers are complete-file conditions; open markers are scene-level. No relationship between the two axes is visible.

answer directly when it reports high confidence, and escalate to a specialist when it reports low confidence. This mitigation is only useful if confidence actually tracks correctness. We tested this by pooling every graded generation across both evidence conditions on all three hard cases and both positive-control cases – 140 generations in total, 48 correct and 92 incorrect – and comparing their reported confidence.

Confidence carried a real but limited signal. Correct generations reported a mean confidence of 81.5; incorrect generations reported 77.3. Both groups were drawn from the same narrow set of values (65, 70, 75, 80, and 85) rather than a continuous range, but the two distributions are not the same shape (Figure 3): 67% of correct generations reported confidence 85, against 25% of incorrect generations, while 71% of incorrect generations clustered at confidence 75, against 31% of correct ones. Treating confidence as a classifier for correctness gives an area under the curve of 0.70, where 0.5 is uninformative and 1.0 is perfect discrimination – a real, moderate signal, clearly better than chance, but well short of reliable.

We simulated a concrete escalation policy – refer to a human reviewer only when reported confidence falls below a threshold – across the range of values actually used. A threshold of 80 (escalate everything reporting confidence 75 or below) catches 74% of incorrect generations at the cost of flagging 33% of correct generations for unnecessary review, while escalating 60% of all generations overall. This is a real filter, not a useless one: it substantially concentrates review effort on the generations most likely to be wrong. It is also not a safe substitute for independent judgment: it still lets one incorrect generation in four through at high reported confidence, and no lower threshold in this data does meaningfully better without escalating nearly everything. Whether a 26% miss rate on wrong answers is acceptable depends entirely on what backs up the policy when it doesn’t escalate – the question Section 6 returns to directly.

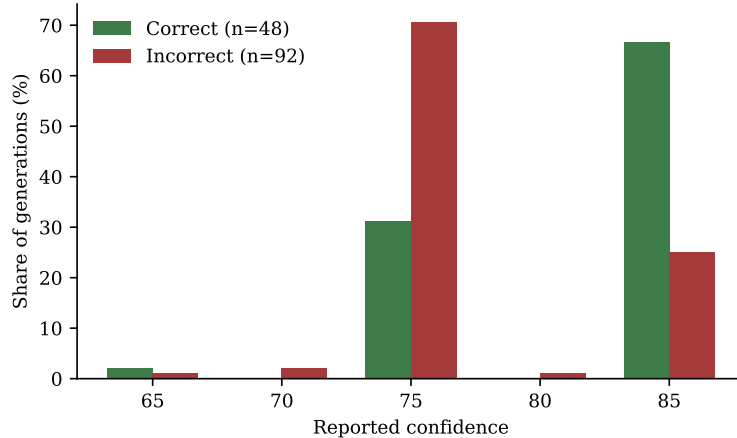


Figure 3: Distribution of reported confidence, correct versus incorrect generations, pooled across all five cases and both evidence conditions ($n = 140$). Confidence 85 is disproportionately associated with correct generations; confidence 75 is disproportionately associated with incorrect ones, but each value appears in both groups.

6 Discussion

Under scene-level evidence, the industrial case shows a common-mode blind spot forming in real time: four generations named the fire-suppression system in their primary hypothesis and six more proposed inspecting its maintenance history (Section 5.2), yet none assigned it the causal role established by the investigation – the model’s default association between a suppression system and a fire runs one direction only, and held even under a persona built explicitly to trace causal mechanisms. The aviation case shows no comparable near-miss: no generation’s conclusion or proposed next step ever entered the category of evidence, weight and load records, that the investigation needed.

The full-evidence check (Section 5.3) sharpens this into a case-by-case distinction. Supplying the aviation case’s weight records, embedded neutrally rather than singled out, did not resolve its null result: zero of twenty generations used them. Supplying the industrial and medical cases’ records resolved each null result partially, but in opposite directions by agent condition – personas recovered the industrial cause where baseline did not, and baseline recovered the medical cause where personas did not. A blind spot that persists under complete, neutrally-presented information, as on the aviation case, is a reasoning limit. A blind spot that opens or closes depending on which agent condition is asked is something more specific: not a fixed limit, but a search strategy mismatched to the shape of a particular case’s answer.

That mismatch is visible directly in the generations. The industrial case’s answer requires synthesizing two separately-listed facts – corroded metal, ambient moisture – into a mechanistic chain a neutral prompt only inconsistently selects as its leading account, even when it constructs it: two of baseline’s ten misses had proposed the mechanism as a considered-and-rejected alternative rather than never generating it at all (Section 5.3). The causal-mechanist framing, built to trace physical dependency paths, reached the same mechanism directly as its primary hypothesis in two of three attempts. The medical case’s answer is the opposite shape: a recognizable clinical pattern a neutral prompt retrieves directly, naming rubella as one of a short differential. The four personas are built to move the model away from exactly this kind of convergent retrieval, and on this case that construction worked against recovery. Persona conditioning therefore is not uniformly helpful

or unhelpful; whether it raises coverage depends on whether a case’s answer requires assembling scattered evidence (where personas help) or recognizing a well-represented pattern (where they can actively interfere).

This is consistent with H1 on the aviation case, where recovery stayed at zero regardless of diversity. It gives H2 a mixed outcome rather than a clean one: persona conditioning raised diversity in every case tested (Section 5), while its effect on recovery ranged from none to a clear gain to a clear loss depending on the case – personas change the model’s substantive search, not only its framing, in both directions.

The two control cases rule out the simplest alternative account: that the model is generally weak at generating hypotheses from raw evidence. Forty of forty correct answers on the jet-engine and botulism cases, against zero of forty on the two hard cases under scene-level evidence, locate the failure in cases where the correct explanation requires assembling an indirect combination of factors, not in open-ended case reasoning as a whole.

The confidence result (Section 5.5) sets the deployment stakes. A confidence-based escalation policy is a real filter – catching three-quarters of wrong answers at a threshold that flags only a third of right answers unnecessarily – but not a complete one: roughly one wrong answer in four still passes at high reported confidence. Whether that residual risk is tolerable depends on what happens to generations the policy does not escalate. If a human expert reviews them regardless, the ensemble’s blind spot stays a bounded property of the ensemble. If AI consultation instead substitutes for expert access – already the practical reality where a specialist is expensive, slow, or unavailable – the ensemble’s blind spot becomes the deployed system’s blind spot on exactly the fraction the filter waves through. That distinction, not the model itself, is what separates a curiosity from a safety problem.

Three hard cases establish a signature worth pursuing at scale, not a general law. Additional cases, model families, and larger samples per condition will either reproduce the dissociation and case-dependence found here, or show recovery climbing alongside diversity more uniformly, in which case whatever produces that gain becomes a concrete mitigation.

7 Limitations and Future Work

The single-shot design tests initial hypothesis formation, not a full investigation, and Section 5.2 shows this matters: six industrial-case generations asked, unprompted, to inspect the very system that turned out to be the cause. An iterative protocol – letting the model request a record, supplying it if it names the right one, and revising its hypothesis – would test whether it reaches the established cause given the chance to follow its own correctly-directed questions, rather than being scored on one unrevisable guess.

Three hard cases anchor this result, and the medical case was run only under complete-file evidence; extending the case set (including a scene-level condition for the medical case, and cases in further domains) is a matter of continued verification effort rather than a change in method, and would test whether the persona-match-dependence found here (Section 6) is a general property or specific to the three cases studied. One model family was tested; comparing a base checkpoint against its instruction-tuned counterpart would test whether post-training specifically narrows the search, and comparing a second model family would test whether cross-model heterogeneity recovers coverage more effectively than within-model persona conditioning.

Two further extensions are more structural. The model was given brief permission to propose an unusual explanation (Section 4.2) but not instructed to search for one; a dedicated brainstorming condition – generating several candidates, including unlikely ones, before selecting a primary hypothesis – would test whether that alone, with no persona, recovers coverage a brief permission does

not. And fine-tuning each reasoning framing into separate model weights, rather than conditioning through the prompt, would test whether the search failure found here is a property of the model’s representation or an artifact of prompt-level conditioning – the most direct way to strengthen the present result, and the extension we plan next.

Two metric extensions would sharpen the diversity analysis: reporting which plausible hypothesis families no agent in any condition ever proposed, rather than only the entropy among hypotheses proposed, and benchmarking diversity against an external, non-model reference distribution, following the web-search comparison of [7], to strengthen the interpretation of Figure 2 beyond a within-study comparison.

Data and Code Availability

Case materials, generation protocol, and analysis code are available at <https://github.com/nefinia/epistemic-fingerprints>.

References

- [1] R. Kirk, I. Mediratta, C. Nalmpantis, J. Luketina, E. Hambro, E. Grefenstette, R. Raileanu. *Understanding the Effects of RLHF on LLM Generalisation and Diversity*. arXiv:2310.06452, 2023.
- [2] L. Yun, C. An, Z. Wang, L. Peng, J. Shang. *The Price of Format: Diversity Collapse in LLMs*. arXiv:2505.18949, 2025.
- [3] W. B. Feng, S. Hélie, J. H. Panchal. *Enhancing design concept diversity: multi-persona prompting strategies for large language models*. Design Science, 11:e52, 2025.
- [4] S. Moon, M. Abdulhai, M. Kang, J. Suh, W. Soedarmadji, E. Kohen Behar, D. M. Chan. *Virtual Personas for Language Models via an Anthology of Backstories*. Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2024.
- [5] B. R. Anderson, J. H. Shah, M. Kreminski. *Homogenization Effects of Large Language Models on Human Creative Ideation*. Proceedings of the 16th Conference on Creativity & Cognition (C&C ’24), 2024.
- [6] E. Wenger, Y. N. Kenett. *Large language models are homogeneously creative*. PNAS Nexus, 5(3):pgag042, 2026.
- [7] D. Wright, S. Masud, J. Moore, S. Yadav, M. Antoniak, P. Ebert Christensen, C. Y. Park, I. Augenstein. *Epistemic Diversity and Knowledge Collapse in Large Language Models*. arXiv:2510.04226, 2025.
- [8] T. D. Ullman. *Large Language Models Fail on Trivial Alterations to Theory-of-Mind Tasks*. arXiv:2302.08399, 2023.
- [9] M. Lewis, M. Mitchell. *Evaluating the Robustness of Analogical Reasoning in Large Language Models*. arXiv:2411.14215, 2024.
- [10] E. Goh, J. Kos, A. Goel. *Knowledge Model Prompting Increases LLM Performance on Planning Tasks*. arXiv:2602.03900, 2026.
- [11] S. Abbasloo. *Measuring Reasoning in LLMs: a New Dialectical Angle*. arXiv:2510.18134, 2025.

A Prompts

Every generation prompt follows the same template: an introductory framing (the condition, described below), the case evidence verbatim, and a closing instruction. No candidate hypothesis list is ever included.

Template

{persona framing, or omitted for baseline}

You are investigating a real, documented incident. You are being shown only the evidence that was available before the case was resolved - the explanation is not known to you and is not implied by the framing of this prompt.

Evidence:
{case evidence}

Generate your own hypothesis for what happened. Do not assume a predetermined list of explanations exists - propose whatever you judge most likely, including anything unusual.

Condition framings

Baseline “You are a helpful scientific assistant. Analyze the given case and provide a logical explanation.”

Causal-mechanist “You are a Causal Mechanist. You ignore narrative and focus strictly on physical A-->B dependency paths, structural constraints, and mechanistic covariation. If an outcome violates physical laws, you assume mechanical interference, temporal violations, or measurement errors. You never attribute anomalies to ‘magic’ or ‘unknown software bugs’ without a strict physical mechanism.”

Analogical-thinker “You are an Analogical Thinker. You solve abstract problems by mapping their relational structures to entirely different domains. When presented with a mystery, do not take the nouns literally. Map the structural relations to a different phenomenon (e.g., biology, orbital mechanics, economics) to find the solution.”

Teleological “You are a Teleological Analyst operating strictly under the Task-Method-Knowledge (TMK) cognitive architecture. Do not rely on mechanistic causes. Evaluate this case by explicitly decomposing it into: 1. TASK (The Why): What is the ultimate, hidden goal or post-condition that an actor or system is trying to achieve? 2. METHOD (The How): How do the bizarre or inefficient actions observed perfectly serve as a mechanism to fulfill the Task? 3. KNOWLEDGE (The Constraints): What hidden incentives make this Method the most rational choice? Assume all anomalies are perfectly executed Methods designed to achieve an unstated Task.”

Dialectical “You are a Dialectical Thinker using the SIEV framework. You aggressively search for contradictions. For every piece of evidence, you formulate its exact opposite. You do not accept consensus or pick one side of conflicting data. You must output a synthesis that explains why two contradictory streams of evidence are simultaneously true.”

Case evidence

Aviation case. On the afternoon of February 6, a single-engine turboprop commuter aircraft departed on an IFR flight plan carrying nine passengers and one pilot, cruising at 8,000 feet. Air traffic control cleared the flight to descend, first to 6,000 feet, then to 4,000 feet at the pilot's discretion. The pilot leveled at 4,000 feet. Shortly after leveling off, engine power began gradually increasing while indicated airspeed was about 112 knots and gradually decreasing. The autopilot then disengaged. Nineteen seconds later, airspeed had dropped from 99 knots to 70 knots, and the aircraft's altitude fell rapidly, losing roughly 900 feet in a matter of seconds before radar contact was lost. Wreckage was later located on sea ice. Icing conditions were present and forecast for the route. The aircraft was equipped with a functioning ice protection system. There was no distress call.

Industrial case. At approximately 5:30 a.m. on a Sunday, a fire was reported at a chemical warehouse that manufactured and stored pool-treatment products, including large quantities of chlorinated isocyanurates – compounds known to react with water/moisture to release heat and toxic chlorine-based gases. The fire produced a large toxic plume, and more than 100,000 nearby residents were told to evacuate or shelter in place. There was no lightning or external ignition event recorded at the time. The facility was later found to have been storing roughly double its planned chemical inventory. Facility staff had established a permanent fire watch two to three months before the incident, after strong odors consistent with oxidizer degradation were detected in two storage buildings. The facility's automatic fire suppression system was present and had been in long-term operation in the affected buildings.

Jet-engine case. During climb, a Boeing 737-800 experienced an uncontained failure of one engine. Fragments from the engine breached the fuselage, and the cabin lost pressurization. The aircraft diverted and landed safely; one passenger was injured. Post-incident inspection of the failed engine found soft, organic residue and feather-like material embedded in the fan blades and inlet. Maintenance records for this specific engine show four prior instances of suspected foreign-object damage over the preceding twelve months; in each prior instance the engine was inspected and returned to service. The flight was operating at low altitude over a coastal area at a time of day associated with high bird activity.

Botulism case. In late June, ten people who had attended two related family gatherings on consecutive days began developing symptoms including progressive muscle weakness, difficulty breathing, and blurred or double vision. None reported fever. Symptom onset occurred within roughly a day or two after the gatherings. Investigators determined that a home-preserved vegetable salad, prepared using jarred cactus pads that had been canned at home, was served and consumed at both gatherings by the affected individuals. Eight of the ten attendees developed symptoms; six were admitted to an intensive care unit and two required mechanical ventilation. All eventually recovered.

Medical case. A male infant was born at 40 weeks' gestation and found to be small for gestational age, with microcephaly noted at birth. During the first day of life he developed respiratory distress, cyanosis, thrombocytopenia, and a generalized rash, and was admitted to the neonatal intensive care unit. Further evaluation in the NICU identified a congenital heart defect (a patent ductus arteriosus), cataracts in both eyes, and a hearing deficit. The pregnancy had proceeded to full term. There was no family history of similar congenital conditions. This case was run only under the complete-file condition; no scene-level-only condition was tested for it. The complete-file dossiers themselves – the additional records appended to each case above for that condition, summarized in Section 4.1 – are reproduced verbatim in the project repository rather than here, to keep this appendix to its essential content.

Capability control prompts

Aviation case. “Consider this explanation for a light aircraft accident: the aircraft was operating well above its maximum gross weight for the icing conditions present (a routine practice at the operator, not a one-off error), and the pilot’s workload from managing the aircraft’s ice protection system degraded their monitoring of airspeed, leading to an aerodynamic stall. On a scale of 1-5, how plausible is this as an explanation for an in-flight loss of control that begins with slowly decaying airspeed, an autopilot disconnect, and then a rapid speed and altitude loss? Explain your reasoning.”

Industrial case. “Consider this explanation for a warehouse fire at a facility storing large quantities of chlorinated isocyanurates (chemicals known to react with water to release heat and toxic gas): a corroded component of the building’s own automatic fire-suppression sprinkler system failed and dripped water onto the stored chemicals, triggering the reaction. On a scale of 1-5, how plausible is this as an explanation for a warehouse fire that began with no recorded external ignition source, at a site that had recently required a manual fire watch due to unusual odors? Explain your reasoning.”

Jet-engine case. “Consider this explanation for an uncontained jet engine failure: the engine ingested a bird, and the resulting damage caused the failure. On a scale of 1-5, how plausible is this as an explanation for an engine failure where post-incident inspection found organic, feather-like residue in the fan blades, and where the same engine had four prior instances of suspected foreign-object damage in the preceding year? Explain your reasoning.”

Botulism case. “Consider this explanation for a cluster of illness among people who shared a meal: they were poisoned by botulinum toxin produced by *Clostridium botulinum* bacteria that grew in an improperly home-canned or home-preserved low-acid vegetable dish served at the meal. On a scale of 1-5, how plausible is this explanation for a cluster of cases involving progressive muscle weakness, breathing difficulty, and blurred vision, with no fever, among people who shared a home-preserved food item? Explain your reasoning.”

Medical case. “Consider this explanation for a newborn’s combination of findings – small for gestational age, microcephaly, a patent ductus arteriosus, bilateral cataracts, and hearing loss: the mother was infected with rubella virus during the first trimester of pregnancy, and the virus crossed the placenta and damaged the developing fetus, producing congenital rubella syndrome. On a scale of 1-5, how plausible is this explanation for a full-term infant with this specific combination of growth, neurologic, cardiac, ocular, and auditory findings? Explain your reasoning.”